

# OpenMP loops: A primer

- OpenMP provides a loop construct that specifies that the iterations of one or more associated loops will be executed in parallel by threads in the team in the context of their implicit tasks.<sup>1</sup>

```
#pragma omp for [clause[ [,] clause] ... ]  
    for (int i=0; i<100; i++) {}
```

- Loop needs to be in canonical form.
- The *clause* can be one or more of the following: `private(...)`, `firstprivate(...)`, `lastprivate(...)`, `linear(...)`, `reduction(...)`, **`schedule(...)`**, `collapse(...)`, `ordered[...]`, `nowait`, `allocate(...)`
- We focus on the clause ***schedule(...)*** in this talk.

1: OpenMP Technical Report 6. November 2017. <http://www.openmp.org/press-release/openmp-tr6/>

# A Schedule for an OpenMP loop

```
#pragma omp parallel for schedule([modifier [modifier]:]kind[,chunk_size])
```

- A ***schedule*** in OpenMP is:
  - a specification of how iterations of associated loops are divided into contiguous non-empty subsets
    - We call each of the contiguous non-empty subsets a *chunk*
  - *and* how these chunks are distributed to threads of the team.<sup>1</sup>
- The *size of a chunk*, denoted as *chunk\_size* must be a positive integer.

1: OpenMP Technical Report 6. November 2017. <http://www.openmp.org/press-release/openmp-tr6/>


# The Kind of a Schedule

- A schedule *kind* is passed to an OpenMP loop schedule clause:
  - provides a hint for how iterations of the corresponding OpenMP loop should be assigned to threads in the team of the OpenMP region surrounding the loop.
- Five *kinds of schedules* for OpenMP loop<sup>1</sup>:
  - *static*
  - *dynamic*
  - *guided*
  - *auto*
  - *runtime*
- The OpenMP implementation and/or runtime defines how to assign chunks to threads of a team given the kind of schedule specified by as a hint.


- › `schedule(static[, chunk_size])`
  - Iterations are divided into **chunks** of `chunk_size`, and chunks are assigned to threads **before entering the loop**
  - If `chunk_size` unspecified, = `NITER/NTHREADS` (with some adjustment...)
  
- › `schedule(dynamic[, chunk_size])`
  - Iterations are divided into chunks of `chunk_size`
  - **At runtime**, each thread requests for a new chunk after finishing one
  - If `chunk_size` unspecified, then = 1

- › `schedule(guided[, chunk_size])`
  - A mix of static and dynamic
  - `chunk_size` determined statically, assignment done dynamically
  
- › `schedule(auto)`
  - Programmer let compiler and/or runtime decide
  - Chunk size, thread mapping..
  - "I wash my hands"
  
- › `schedule(runtime)`
  - Only runtime decides according to run-sched-var ICV
  - If run-sched-var = `auto`, then implementation defined


```
a();  
#pragma omp parallel for  
for (int i = 0; i < 16; ++i) {  
    w(i);  
}  
z();
```



```
a();  
#pragma omp parallel for schedule(static,1)  
for (int i = 0; i < 16; ++i) {  
    w(i);  
}  
z();
```

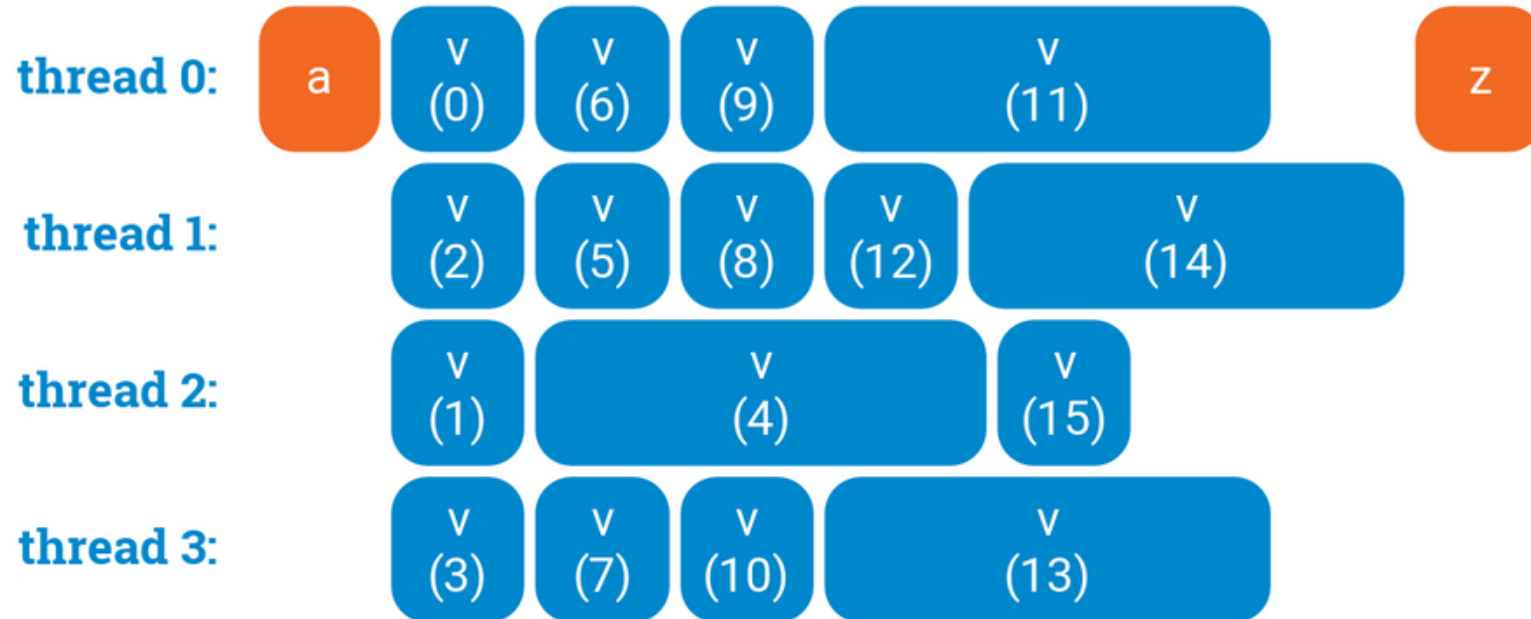


```
a();  
#pragma omp parallel for schedule(static,2)  
for (int i = 0; i < 16; ++i) {  
    w(i);  
}  
z();
```





```
a();  
#pragma omp parallel for schedule(dynamic,1)  
for (int i = 0; i < 16; ++i) {  
    v(i);  
}  
z();
```



## **Static vs. Dynamic**

- Static has less initialization cost
- Dynamic has better load balancing
- Dynamic scheduling is expensive (setup, fetching work, loop exit)