


```
V=alpha();  
W=beta();  
X=gamma(v,w);  
Y=delta();  
printf("%g\n", epsilon(x,y));
```



Data dependence diagram

Functions alpha, beta, delta may be executed in parallel

Worksharing **sections** Directive

sections directive enables specification of task parallelism

- Sections construct gives a different structured block to each thread.

```
#pragma omp sections [clause list]
    private (list)
    firstprivate (list)
    lastprivate (list)
    reduction (operator: list)
    nowait
```

```
{
#pragma omp section
    structured_block
#pragma omp section
    structured_block
}
```

```

#include "omp.h"
#define N 1000
int main(){
    int i;
    double a[N], b[N], c[N], d[N];
    for(i=0; i<N; i++){
        a[i] = i*2.0;
        b[i] = i + a[i]*22.5;
    }
    #pragma omp parallel shared(a,b,c,d) private(i)
    {
        #pragma omp sections nowait
        {
            #pragma omp section
            for(i=0; i<N; i++) c[i] = a[i]+b[i];
            #pragma omp section
            for(i=0; i<N; i++) d[i] = a[i]*b[i];
        }
    }
}

```

Two tasks are
computed
concurrently

By default, there is a barrier at the end of the sections. Use the "nowait" clause to turn off the barrier.


```
#include "omp.h"

#pragma omp parallel
{
    #pragma omp sections
    {
        #pragma omp section
        v=alpha();
        #pragma omp section
        w=beta();
    }
    #pragma omp sections
    {
        #pragma omp section
        x=gamma(v,w);
        #pragma omp section
        y=delta();
    }
    printf("%g\n", epsilon(x,y));
}
```

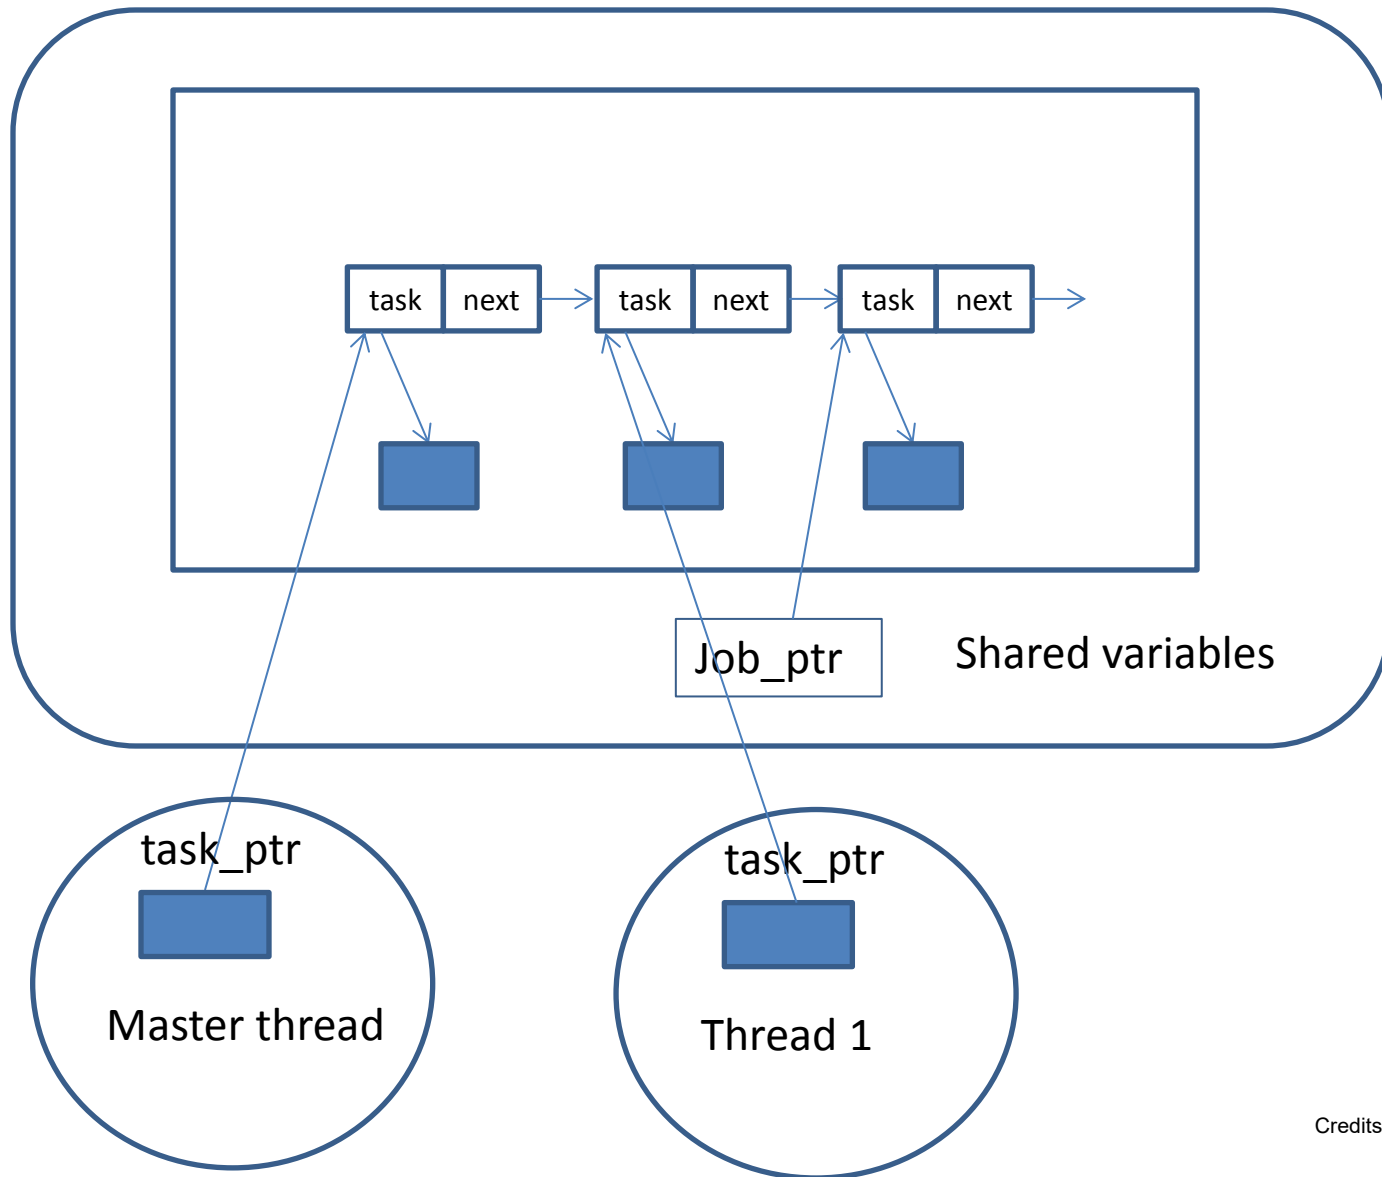
Code Fragment for Manager/Worker Model

```
int main(int argc, char argv[])
{
    struct job_struct job_ptr;
    struct task_struct *task_ptr;
    ...
    task_ptr = get_next_task(&job_ptr);
    while(task_ptr != NULL){
        complete_task(task_ptr);
        task_ptr = get_next_task(&job_ptr);
    }
    ...
}

struct task_struct *get_next_task(struct job_struct *job_ptr)
{
    struct task_struct *answer;
    if(job_ptr == NULL) answer = NULL;
    else
    {
        answer = job_ptr->task;
        job_ptr = job_ptr->next;
    }
    return answer;
}
```



- Two threads complete the work



Tasking

OpenMP 3.0 and Tasks

Tasks allow to parallelize irregular problems

- Unbounded loops
- Recursive algorithms
- Manager/work schemes
- ...

A task has

- **Code** to execute
- **Data** environment (It owns its data)
- **Internal control variables**
- An assigned thread that executes the code and the data

Two activities: packaging and execution

- Each encountering thread packages a new instance of a task (code and data)
- Some thread in the team executes the task at some later time

Recursive approach to compute Fibonacci



```
int main(int argc,  
        char* argv[])  
{  
    [...]  
    fib(input);  
    [...]  
}
```

```
int fib(int n)    {  
    if (n < 2) return n;  
    int x = fib(n - 1);  
    int y = fib(n - 2);  
    return x+y;  
}
```

- On the following slides we will discuss three approaches to parallelize this recursive code with Tasking.

C/C++

```
#pragma omp task [clause]  
... structured block ...
```

Fortran

```
!$omp task [clause]  
... structured block ...  
!$omp end task
```

■ Each encountering thread/task creates a new Task

- Code and data is being packaged up
- Tasks can be nested
 - Into another Task directive
 - Into a Worksharing construct

■ Data scoping clauses:

- `shared(list)`
- `private(list)` `firstprivate(list)`
- `default(shared | none)`

- **Some rules from *Parallel Regions* apply:**
 - Static and Global variables are shared
 - Automatic Storage (local) variables are private

- **If shared scoping is not derived by default:**
 - Orphaned Task variables are `firstprivate` by default!
 - Non-Orphaned Task variables inherit the `shared` attribute!
 - Variables are `firstprivate` unless `shared` in the enclosing context

First version parallelized with Tasking (omp-v1)




```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```


```
int fib(int n)    {
    if (n < 2) return n;
    int x, y;
    #pragma omp task shared(x)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

- **Only one Task / Thread enters `fib()` from `main()`, it is responsible for creating the two initial work tasks**
- **Taskwait is required, as otherwise `x` and `y` would be lost**


- T1 enters fib(4)
- T1 creates tasks for fib(3) and fib(2)
- T1 and T2 execute tasks from the queue
- T1 and T2 create 4 new tasks
- T1 - T4 execute tasks




Task Queue



- T1 enters fib(4)
- T1 creates tasks for fib(3) and fib(2)
- T1 and T2 execute tasks from the queue
- T1 and T2 create 4 new tasks
- T1 - T4 execute tasks
- ...



■ Overhead of task creation prevents better scalability!



- If the expression of an **if** clause on a task evaluates to **false**
 - The encountering task is suspended
 - The new task is executed immediately
 - The parent task resumes when the new task finishes
 - Used for optimization, e.g., avoid creation of small tasks


■ Improvement: Don't create yet another task once a certain (small enough) n is reached

```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```

```
int fib(int n)    {
    if (n < 2) return n;
    int x, y;
    #pragma omp task shared(x) \
        if(n > 30)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y) \
        if(n > 30)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

- Speedup is ok, but we still have some overhead when running with 4 or 8 threads

Speedup of Fibonacci with Tasks




- **Improvement: Skip the OpenMP overhead once a certain n is reached (no issue w/ production compilers)**

```
int main(int argc,
        char* argv[])
{
    [...]
    #pragma omp parallel
    {
        #pragma omp single
        {
            fib(input);
        }
    }
    [...]
}
```

```
int fib(int n)    {
    if (n < 2) return n;
    if (n <= 30)
        return serfib(n);
    int x, y;
    #pragma omp task shared(x)
    {
        x = fib(n - 1);
    }
    #pragma omp task shared(y)
    {
        y = fib(n - 2);
    }
    #pragma omp taskwait
    return x+y;
}
```

■ Everything ok now 😊

Speedup of Fibonacci with Tasks



Data Scoping Example (1/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a:
            // Scope of b:
            // Scope of c:
            // Scope of d:
            // Scope of e:

        }
    }
}
```

Data Scoping Example (2/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b:
            // Scope of c:
            // Scope of d:
            // Scope of e:

        }
    }
}
```

Data Scoping Example (3/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c:
            // Scope of d:
            // Scope of e:
        }
    }
}
```

Data Scoping Example (4/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d:
            // Scope of e:

        }
    }
}
```


Data Scoping Example (5/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d: firstprivate
            // Scope of e:

        }
    }
}
```

Data Scoping Example (6/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared
            // Scope of b: firstprivate
            // Scope of c: shared
            // Scope of d: firstprivate
            // Scope of e: private
        }
    }
}
```

Hint: Use default(none) to be forced to think about every variable if you do not see clear.

Data Scoping Example (7/7)



```
int a = 1;
void foo()
{
    int b = 2, c = 3;
    #pragma omp parallel shared(b)
    #pragma omp parallel private(b)
    {
        int d = 4;
        #pragma omp task
        {
            int e = 5;

            // Scope of a: shared,          value of a: 1
            // Scope of b: firstprivate,    value of b: 0 / undefined
            // Scope of c: shared,          value of c: 3
            // Scope of d: firstprivate,    value of d: 4
            // Scope of e: private,        value of e: 5

        }
    }
}
```

■ OpenMP `barrier` (implicit or explicit)

- All tasks created by any thread of the current *Team* are guaranteed to be completed at barrier exit

C/C++

```
#pragma omp barrier
```

■ Task barrier: `taskwait`

- Encountering Task suspends until child tasks are complete
 - Only direct childs, not descendants!

C/C++

```
#pragma omp taskwait
```

■ Task Synchronization explained:

```
#pragma omp parallel num_threads(np)
{
    #pragma omp task
        function_A();
    #pragma omp barrier
    #pragma omp single
    {
        #pragma omp task
            function_B();
    }
}
```

np Tasks created here, one for each thread

All Tasks guaranteed to be completed here

1 Task created here

B-Task guaranteed to be completed here

Loop Collapse

- Allows parallelization of perfectly nested loops without using nested parallelism
- Compiler forms a single loop and then parallelizes this

```
{  
  ...  
  #pragma omp parallel for collapse (2)  
  for(i=0;i< N; i++)  
  {  
    for(j=0;j< M; j++)  
    {  
      foo(A,i,j);  
    }  
  }  
}
```