

A quick revision

Aravind Sukumaran Rajam

What happened?

- Travel Issues
- Thank fully Prof. Ananth stepped in



What are we going to do today?

- Why is this course important?
- How will this help us in our future (academic / industry)?
- My vision for this course
- What I promise
- What I want you to promise me
- A quick overview of the syllabus
- A quick revision of what you have learned
- Data dependencies
- Introduction to OpenMP

Why is this course important?


- There is a saying that “Mathematics is the language of science”
- Today, we can tell that “High Performance Computing (HPC)” is the beating heart of computer science, along with machine learning and computer security
 - Parallel computing is at the core of HPC
- Nothing is sequential anymore... Why? (we will see soon)
- Parallel computing knowledge is a basic requirement now a days
- Machine learning, scientific computing, video games, computer simulations, all relies on parallel computing

Why is this course important?

- Tremendous opportunities in the industry
- Knowing parallel computing will help you a lot in academic and industry research
 - BTW you should do a PhD or Masters (more on this in the future)
 - Want to try HPC? This course is page 1 😊
- You will definitely become a better programmer
- Understand the hardware better
- Improve algorithmic knowledge

Gems!!!

- This is what will happen if you don't have parallel computing knowledge



Credits: To the unknown creator

My vision for this course

- My overall vision for this course is to make sure that you **understand** the basics of Parallel Computing
 - There is a big difference between scoring in an exam and understanding something
- To get you interested in HPC
- To get you interested in HPC related Graduate Studies

What I promise

- I will take every effort to make sure that you **understand** what I teach
- I will actively encourage you to think openly
- I will make every effort to accommodate your feedbacks into the course
 - Not next year!!! This year itself!!!
 - If I can't accommodate a request, I will tell you the reason
 - You will be provided plenty of opportunities to provide anonymous feedbacks

What I want you to promise me

- After every lecture, the same day you will at least spent 15 minutes to think about what we have learned

What I want you to promise me



Doing this is not going to help you!!!

Credits: To the unknown creator

What I want you to promise me

- After every lecture, the same day you will at least spent 15 minutes to think about what we have learned
- You will not copy assignments
- You will be active in the class and ask a lot of questions
 - There are no stupid questions
- You will actively give feedbacks to me
- You will put every effort to succeed in this course

A quick overview of the syllabus

```
int val = 3 / 2 * 4;  
printf("%d", val);
```

- a) 6
- b) 0.375
- c) 4
- d) NA

Do you know how
many people have
made this mistake

A quick overview of the syllabus

- Introduction to fundamental parallel computing concepts
- A quick review of computer architecture (parallel units and caches)
- Data dependence analysis
- Introduction to shared memory parallelization:
 - Introduction to OpenMP
 - Task level parallelization
 - Loop level parallelization
- Advanced OpenMP
 - Parallel algorithms
 - Challenges associated with sparse computations
- Performance monitoring and debugging
- Introduction to distributed computing
- Introduction to MPI programming
- Advanced MPI
 - Visualization and performance monitoring
 - Advanced MPI API's
 - Challenges associated with MPI
- To infinity and beyond
 - What are the limitations of OpenMP and MPI
 - Can we design better parallel protocols
 - Future courses

A quick revision of what you have learned


- Why parallel computing?
 - Why not go with sequential computing

End of Moore's Law

Problem:

- How to achieve good **performance**?
- Increase clock frequency?
 - stalled around 2000*
- Transistor count?
 - Nearing physical limits
 - 5nm? (maybe 3nm)

Stuttering



A quick revision of what you have learned

- Why parallel computing?
 - Why not go with sequential computing
 - What are the advantages?
 - What are the disadvantages?



Advantages?




Disdvantages?

A quick revision of what you have learned

- Name all the parallel processing hardware units (starting from highest level)
 - Distributed nodes
 - Shared memory (multi-core)
 - Simultaneous multithreading (SMT)
 - Single Instruction Multiple Data (SIMD) – Vector units
 - Pipeline parallelism


Parallelism vs Concurrency

- **Concurrency:** The process by which multiple jobs can be executed in an overlapped manner, one at a time




Parallelism vs Concurrency

- **Concurrency:** The process by which multiple jobs can be executed in an overlapped manner, one at a time




Parallelism vs Concurrency

- Parallelism: The process by which multiple jobs can be executed at the same time



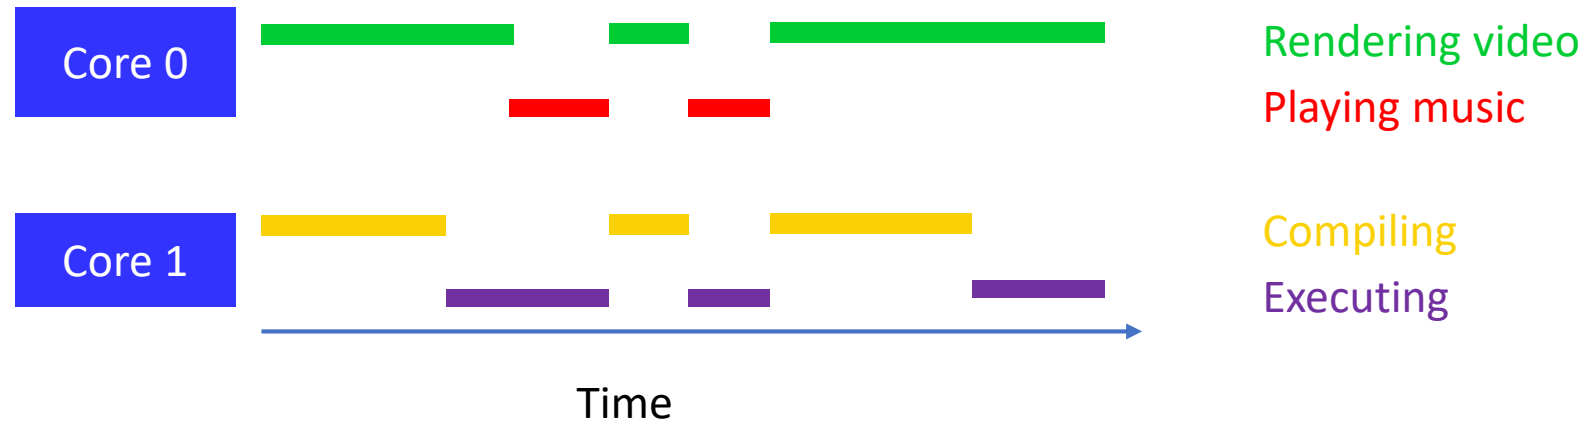
Parallelism vs Concurrency

- Parallelism: The process by which multiple jobs can be executed at the same time



Parallelism vs Concurrency

- Can we have both parallelism and concurrency?



Which of this is false?

- a) An application can be concurrent and parallel
- b) An application can be concurrent and not parallel
- c) An application can be not concurrent and parallel
- d) An application can be not concurrent and not parallel

And the answer is None 😊

Speedup

- What is speedup
 - $SUP(P,N) = T_{seq(N)} / T_{parallel(N, P)}$
 - $T_{seq(N)}$ -- time for sequential processing
 - $T_{parallel(N, P)}$ -- time for parallel processing on P processors
 - N – problem size
 - How can we quantify?

Efficiency

- What is efficiency
 - $EFF(P,N) = T_{seq(N)} / (P * T_{parallel(N, P)}) = SUP(P,N) / P$
 - $T_{seq(N)}$ -- time for sequential processing
 - $T_{parallel(N, P)}$ -- time for parallel processing on P processors
 - N – problem size
 - What is the intuitive meaning?
 - What do you expect $EFF(P,N)$ to be?
 - What is the best case?
 - Can it ever happen: $EFF(P,N) > 1$?
 - What is the worst case?
 - Can it ever happen: $T_{parallel(N, P)} > T_{seq(N)}$
 - Imagine that $EFF(P,N) = 0.3$
 - Is this is a good thing?
 - Is this is a bad thing?

Scalability

- What is scalability
 - The efficiency as a function of number of parallel units
- Wait wait wait
 - Efficiency???
 - a) Should we measure the efficiency for a fixed problem size and a variable number of parallel units?
 - b) Should we measure the efficiency by increasing both problem size and number of parallel units?

Scalability

- Measure the efficiency for a fixed problem size and a variable number of parallel units?
 - What is this question intuitively trying to answer?
 - Strong scaling (how solution time varies for a fixed problem size as we add more processors)
 - Related to Amdahl's law:
 - Let $t_{\text{seq_proportion}}$ be the proportion of time required to execute the sequential part
 - Let $t_{\text{par_proportion}}$ be the proportion of time required to execute the parallelizable part
 - **Strong scaling speedup** = $1 / (t_{\text{seq_proportion}} + t_{\text{par_proportion}} / P)$
 - What happens as P increases?
 - What does this imply?
 - For a fixed problem size, the upper bound of speedup is limited by serial part of a program.
 - Let's try to draw a graph (speedup vs P for $t_{\text{seq_proportion}}$)

Scalability

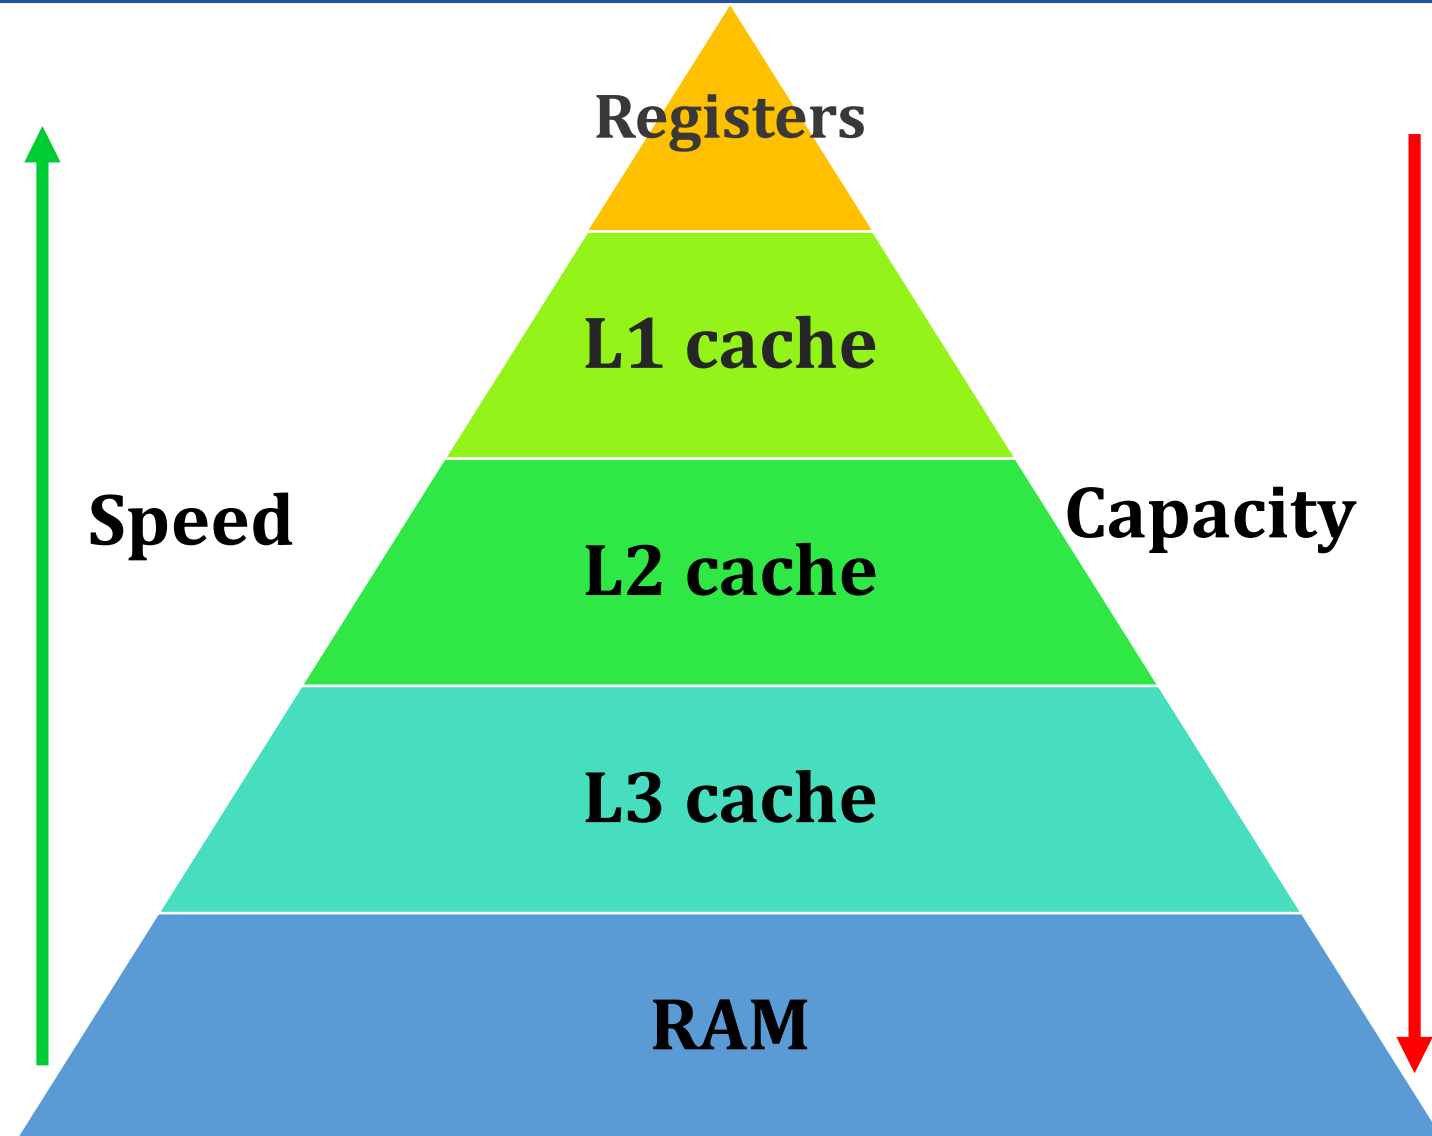
- Should we measure the efficiency by increasing both problem size and number of parallel units?
 - What is this question intuitively trying to answer?
 - Weak scaling (how the solution time varies when the problem size increases, but the work per parallel unit is kept constant)
 - Related to Gustafson's law:
 - In many cases,
 - the sequential execution time is independent of the problem size
 - The parallel part scales linearly with the number of processors (example?)
 - In such cases the weak scaling speedup can be written as $(t_{\text{seq}}_{\text{proportion}} + t_{\text{par}}_{\text{proportion}} * P)$

Scalability

- What are the factors that limit parallel scalability
 - Communication overheads
 - Startup costs
- In general Amdahl's law is more important and difficult to achieve

CPU Memory hierarchy

- **Registers:**
 - Capacity: ~32
 - Access time: ~1 clock cycle
- **L1 cache:**
 - Capacity: ~32KB
 - Access time: ~4-5 clock cycles
- **L2 cache:**
 - Capacity: ~256KB
 - Access time: ~10 clock cycles
- **L3 cache:**
 - Capacity: ~4 to 35MB
 - Access time: ~65 clock cycles
- **Ram:** ~300 clock cycles



CPU Memory hierarchy

- Keep the data as close to the processor as possible
 - How to do this???
 - Unfortunately formalization of this process is out of context for this course (The compiler course covers this)
 - However, we can build an intuition based on data dependencies

Data dependencies

- Two statements are said to be data-dependent if both of them accesses the same memory location **at least one of them is a write operation**
- Constraints the order in which the statements are scheduled

```
int a = 10;  
int b = 20;  
int c = a+b;  
printf("%d",c);
```

Original Code

```
int b = 20;  
int a = 10;  
int c = a+b;  
printf("%d",c);
```

Valid schedule

```
int a = 10;  
int b = 20;  
printf("%d",c);  
int c = a+b;
```

Invalid schedule

Flow dependence

S1: $a = b + c$
S2: $d = a + 42$
S3: $a = 7;$
S4: $e = a + d$

- Also called as Read-After-Write (RAW) dependence
- S_t (target statement) reads a memory that was written by S_s (source statement)
- E.g.
 - $S1 \rightarrow S2$ on 'a'
 - $S2 \rightarrow S4$ on 'd'

Anti dependence

S1: $a = b + c$
S2: $d = a + 42$
S3: $a = 7;$
S4: $e = a + d$

- Also called as Write-After-Read (WAR) dependence
- S_s (source statement) reads a memory location that will be overwritten by S_t (target statement)
- E.g.
 - $S2 \rightarrow S3$ on 'a'
- Can you think about a way to avoid this dependence?

Output dependence

S1: $a = b + c$
S2: $d = a + 42$
S3: $a = 7;$
S4: $e = a + d$

- Also called as Write-After-Write (WAW) dependence
- S_s (source statement) writes to a memory location that will be overwritten by S_t (target statement)
- E.g.
 - $S1 \rightarrow S3$ on 'a'
- Can you think about a way to avoid this dependence?

Input dependence

S1: $a = b + c$

S2: $d = a + 42$

S3: $e = b * c$

- Also called as Read-After-Read (RAR) dependence
- S_s (source statement) reads a memory location that will be read by S_t (target statement)
- E.g.
 - S1 \rightarrow S3 on 'b' and 'c'
- Is this really a dependence
- Should we ever consider this?

Data dependences in Loops

```
for (int i = 1; i < 10; i++)  
    A[i] = B[i-1] + C[i]    // S1  
    B[i] = A[i+2] * C[i]    // S2
```

- What are all the dependences?
- The dependencies are based on dynamic instances

Data dependences in Loops

```
for (int i = 1; i < 10; i++)  
    A[i] = B[i-1] + C[i]    // S1  
    B[i] = A[i+2] * C[i]    // S2
```

- One easy way to see the dependences in loops is to unroll them

```
i=1 { A[1] = B[0] + C[0]  
      B[1] = A[3] * C[0]  
i=2 { A[2] = B[1] + C[1]  
      B[2] = A[4] * C[1]  
i=3 { A[3] = B[2] + C[2]  
      B[3] = A[5] * C[2]
```

```
A[i] -> A[i+2] //WAR(from S2 to S1)  
B[i] -> B[i+1] //RAW(from S2 to S1)  
C[i] -> C[i]   //RAR
```

Data dependences

- Two statements are said to be data-dependent if both of them accesses the same memory location **at least one of them is a write operation**
- Constrains the order in which the statements are scheduled
- **Why would we ever deviate from the order in which user wrote the code:**
 - Cache friendliness
 - Allows parallelization

Data dependences

- For a loop to be parallel, it should not carry any data dependences

```
for (int i = 1; i < 10; i++)  
  A[i] = B[i-1] + C[i]    // S1  
  B[i] = A[i+2] * C[i]    // S2
```

i loop carries dependence;
cant be parallelized

```
for (int i = 1; i < 10; i++)  
  A[i] = B[i-1] + C[i]    // S1  
  D[i] = E[i+2] * C[i]    // S2
```

i loop does not carry dependence;
can be parallelized

Data dependences

- For a loop to be parallel, it should not carry any data dependences

```
for (int i = 1; i < 10; i++)  
  for(int j = 1; j < 10; j++)  
    A[i][j] = A[i][j-1] + C[i]
```

i loop does not carry dependence;
j loop carries dependence;
i loop **can** be parallelized
j loop **cant** be parallelized

```
for (int i = 1; i < 10; i++)  
  for(int j = 1; j < 10; j++)  
    A[i][j] = A[i+1][j] + C[i]
```

i loop carries dependence;
j loop does not carry dependence;
i loop **cant** be parallelized
j loop **can** be parallelized

Data dependences

- For a loop to be parallel, it should not carry any data dependences
- By now, you know how to check if a loop can be parallelized or not manually. But how do you automate this?
 - Take the compiler class next semester 😊
- A lot of times, we concentrate on loop level parallelism. Why?

Data dependences

- There are loop transformations that can enable parallelism on a given loopnest (or a part of it)
 - For details take the compiler class 😊
- A simple example is given below (Loop Fission)

```
for (int i = 2; i < 10; i++)  
    A[i] = B[i+1] + C[i]    // S1  
    B[i] = A[i-2] * C[i]    // S2
```

i loop carries dependence;
cant be parallelized

```
for (int p = 2; p < 10; p++)  
    A[p] = B[p+1] + C[p]    // S1  
for (int q = 2; q < 10; q++)  
    B[q] = A[q-2] * C[q]    // S2
```

p and q does not carry dependence;
Both **can** be parallelized

- Pei Yau (pei-yau.weng@wsu.edu)