

CS 598: Communication Cost Analysis of Algorithms

Lecture 1: Course motivation and overview; collective communication

Edgar Solomonik

University of Illinois at Urbana-Champaign

August 22, 2016

A simple model for point-to-point messages

The time to send or receive a message of s bytes is

$$T_{\text{sr}}^{\alpha,\beta}(s) = \alpha + s \cdot \beta$$

- α – **latency/synchronization cost** per message
- β – **bandwidth cost** per byte
- each processor can send and/or receive one message at a time

Let P processors send a message of size s in a ring,

- the **communication volume** (total amount of data sent) is $P \cdot s$
- What is the **communication cost** (α - β -model execution time)?
 - if the messages are sent *simultaneously*,

$$T_{\text{sim-ring}}^{\alpha,\beta}(s) = T_{\text{sr}}^{\alpha,\beta}(s) = \alpha + s \cdot \beta$$

- if the messages are sent *in sequence*,

$$T_{\text{seq-ring}}^{\alpha,\beta}(s, P) = P \cdot T_{\text{sr}}^{\alpha,\beta}(s) = P \cdot (\alpha + s \cdot \beta)$$

Broadcasts in the α - β model

The execution time of a broadcast of a message of size s to P processors is

- using a **binary tree** of height


$$h_r = 2(\log_2(P+1) - 1) \approx 2\log_2(P)$$

$$\begin{aligned} T_{\text{bcast-bnr}}^{\alpha,\beta}(s, P) &= h_r \cdot T_{\text{sr}}^{\alpha,\beta}(s) \\ &= h_r \cdot (\alpha + s \cdot \beta) \end{aligned}$$

- using a **binomial tree** of height

$$h_m = \log_2(P+1) \approx \log_2(P)$$

$$\begin{aligned} T_{\text{bcast-bnm}}^{\alpha,\beta}(s, P) &= h_m \cdot T_{\text{sr}}^{\alpha,\beta}(s) \\ &= h_m \cdot (\alpha + s \cdot \beta) \end{aligned}$$



Therefore, a binomial tree broadcast is $h_r/h_m \approx 2$ faster than a binary tree broadcast in the α - β model

Large-message broadcasts

Lets now consider broadcasts of a message of a size $s \geq P$ bytes

- recall binomial tree broadcast cost:

$$T_{\text{bcast-bnm}}^{\alpha-\beta}(s, P) = \log_2(P+1) \cdot (\alpha + s \cdot \beta)$$

- consider instead the following broadcast schedule
 - the root sends a different segment of the message to each processor
 - all processors exchange segments in $P - 1$ near-neighbor ring exchanges
- the cost of this broadcast schedule is

$$\begin{aligned} T_{\text{bcast-ring}}^{\alpha-\beta}(s, P) &= (P - 1)(T_{\text{sr}}^{\alpha-\beta}(s/P) + T_{\text{sim-ring}}^{\alpha-\beta}(s/P)) \\ &= 2(P - 1)(\alpha + s/P \cdot \beta) \approx 2(P \cdot \alpha + s \cdot \beta) \end{aligned}$$

- for sufficiently large message sizes, the new schedule is faster,

$$\lim_{s \rightarrow \infty} \left(\frac{T_{\text{bcast-bnm}}^{\alpha-\beta}(s, P)}{T_{\text{bcast-ring}}^{\alpha-\beta}(s, P)} \right) \approx \log_2(P)/2$$

Pipelined binary tree broadcast

Send a packet of size k to left child then to right child

- as before, total message size s , tree height $h \approx \log_2(P)$
- each message costs $\alpha + k \cdot \beta$
- root sends $2s/k$ messages
- last packet takes $2h$ sends to reach rightmost tree leaf
- therefore, the total cost expression is

$$\begin{aligned} T_{\text{PBT}}^{\alpha,\beta}(s, P, k) &\approx 2(h + s/k)(\alpha + k \cdot \beta) \\ &= 2(h \cdot \alpha + s \cdot \beta + (s/k) \cdot \alpha + hk \cdot \beta) \end{aligned}$$

- we can now derive the optimal message size

$$k_{\text{opt}}^{\alpha,\beta}(s, P) = \underset{k}{\operatorname{argmin}}(T_{\text{PBT}}^{\alpha,\beta}(s, P, k)) = \sqrt{\frac{s \cdot \alpha}{h \cdot \beta}}$$

- furthermore, $T_{\text{PBT}}^{\alpha,\beta}(s, P, k_{\text{opt}}^{\alpha,\beta}(s, P)) \approx 2(\sqrt{h \cdot \alpha} + \sqrt{s \cdot \beta})^2$,
a factor of 2 more expensive than the Träff and Ripke protocol

Double Tree

Observation: *the leaves of a binary tree, $(P - 1)/2$ processors, send nothing, while the internal nodes do all the work.*

Double Pipelined Binary Tree Broadcast

- define two pipelined binary trees with a shared root
- non-root processors act as a leaf in one and as an internal node in the second
- send half of the message down each tree, alternating directions with packets of size k




Diagram taken from: Hoefer, Torsten, and Dmitry Moor. "Energy, Memory, and Runtime Tradeoffs for Implementing Collective Communication Operations."

Double pipelined binary tree

The cost of the double pipelined binary tree is essentially the same as the cost of a single pipelined binary tree with half the message size,

$$T_{\text{DPBT}}^{\alpha, \beta}(s, P) \approx 2h \cdot \alpha + 2\sqrt{2s \cdot h} \cdot \sqrt{\alpha \cdot \beta} + \textcolor{red}{s} \cdot \textcolor{red}{\beta}$$

for a sufficiently large message size (s) this is twice as fast as a single pipelined binary tree.

Other types of collective communication

We can classify collectives into four categories

- **One-to-All:** Broadcast, Scatter
- **All-to-One:** Reduce, Gather
- **All-to-One + One-to-All:** Allreduce (Reduce+Broadcast), Allgather (Gather+Broadcast), Reduce-Scatter (Reduce+Scatter), Scan
- **All-to-All:** All-to-all

MPI (Message-Passing Interface) provides all of these as well as variable size versions (e.g. (All)Gatherv, All-to-allv), see online for specification of each routine.

We now present protocols for these and their cost in the $\alpha - \beta$ model, with

$$s = \begin{cases} \text{input size} & : \text{one-to-all collectives} \\ \text{output size} & : \text{all-to-one collectives} \\ \text{per-processor input/output size} & : \text{all-to-all collectives} \end{cases}$$

Tree collectives

We have demonstrated how (double/pipelined) binary trees and binomial trees can be used for broadcasts

- *A reduction may be done via any broadcast tree with the same communication cost, with reverse data flow*

$$T_{\text{reduce}} = T_{\text{broadcast}} + \text{cost of local reduction work}$$

Scatter is strictly easier than broadcast, pipeline half message to each child in a binary tree

$$T_{\text{scatter}}^{\alpha, \beta}(s, P) \approx 2 \log_2(P) \cdot \alpha + s \cdot \beta$$

- *A gather may be done via the reverse of any scatter protocol:*

$$T_{\text{gather}} = T_{\text{scatter}}$$

All-to-One + One-to-All collectives can be done via two trees, but is this most efficient? What about **All-to-All** collectives?